

Red Hat SRE Challenge: Dockerfile Sources

The challenge

As part of our compliance requirements, we want to make sure that we are building containers with trusted source images.

Your task is to build a tool that given a list of repositories, it identifies all the `Dockerfile` files inside each repository, extracts the image names from the `FROM` statement, and returns a `json` with the aggregated information for all the repositories.

You can find the details of the `FROM` command here:

<https://docs.docker.com/engine/reference/builder/#from>

The input will be provided as a URL pointing to a plaintext file. Each line will have two fields separated by a space:

- the https url of the github public repository
- the commit SHA to verify.

You can skip any line that doesn't match this pattern.

Example input:

<https://gist.githubusercontent.com/jmelis/c60e61a893248244dc4fa12b946585c4/raw/25d39f67f2405330a6314cad64fac423a171162c/sources.txt>

Example output:

```
{
  "data": {
    "https://github.com/app-sre/qontract-reconcile.git:30af65af14a2dce962df923446afff24dd8f123e": {
      "dockerfiles/Dockerfile": [
        "quay.io/app-sre/qontract-reconcile-base:0.2.1"
      ]
    },
    "https://github.com/app-sre/container-images.git:c260deaf135fc0efaab365ea234a5b86b3ead404": {
      "jiralert/Dockerfile": [
        "registry.access.redhat.com/ubi8/go-toolset:latest",
        "registry.access.redhat.com/ubi8-minimal:8.2"
      ],
      "qontract-reconcile-base/Dockerfile": [
        "registry.access.redhat.com/ubi8/ubi:8.2",
        "registry.access.redhat.com/ubi8/ubi:8.2",
        "registry.access.redhat.com/ubi8/ubi:8.2"
      ]
    }
  }
}
```

Things to take into account

- We expect the exercise to take around 4 hours. It could be more or less depending on your experience.
- We would like the codebase to be in either of the two languages most commonly used by our team: Python or Golang.
- Do not rush the exercise. Time to delivery will not be part of the evaluation criteria. We care about the quality of the code and making reasoned choices.
- **Execution time will be part of the evaluation criteria.** Considering n as the number of repositories to be processed, we expect an implementation with better time complexity than $O(n)$.
- Invalid lines in the input file should be ignored.
- This should be production grade code.
- You can extend the return payload to include errors.

Deliverables

- URL to *private* GitHub repository with the code and the assignment PDF.
- A README.md file detailing your implementation and any additional features added.
- The Hiring Specialist will share the GitHub username of the reviewer with you. You will need to add them to the private repository with read access.

Bonus points

- Since we are a cloud-native team, we want to run this as a [Kubernetes Job](#). If you already know kubernetes, that is excellent. If you don't, we will hugely value you taking the time to check out [minikube](#) and figuring out how to use `Jobs`. The list of repository urls should be provided to the Job with the `REPOSITORY_LIST_URL` environment variable, which should point at an url.
- Please feel free to implement any additional features that make this project more production ready. Do make sure to document them in the README.